

Original code run time vs final code (mean over 10 runs on stan1):

Original: 47.14s Final: 1.65s

Arithmetic and repeated calculation:

“/pow(dx,2.0)” is expensive as it involves division and calling pow, set a variable $dx2r = 1/(dx*dx)$

Similarly, “/pow(dy,2.0)” $\rightarrow dy2r = 1/(dy*dy)$, “/dx/2.0” $\rightarrow dxr = 0.5/dx$, “/dy/2.0” $\rightarrow dyr = 0.5/dy$

“dt*nu” $\rightarrow dtnu = dt*nu$, this is a repeated multiplication, saves little but nonzero amount of time

“Lapl” and “grad” don’t need to be set to 0.0 every time they are used, in the ix pass, they can just be set equal to their calculation rather than added.

These constants are defined outside of the istep loop to make sure they aren’t repeatedly redefined.

Factorising:

Factorising reduces the amount of multiplication and other arithmetic.

The current general form of Lapl, grad and u_new are (i,j depend on the condition, d2r and dr are $dx2r/dy2r$ and dxx/dyy):

Lapl: $-2.0*u[i1][j1]*d2r + u[i2][j2]*d2r + u[i3][j3]*d2r$

grad: $(u[i4][j4] - u[i5][j5])*dr$

u_new: $u[ix][iy] - dt*u[ix][iy]*grad + dtnu*Lapl$

They can be factorised to:

Lapl: $(-2.0*u[i1][j1] + u[i2][j2] + u[i3][j3])*d2r$

grad: $(u[i4][j4] - u[i5][j5])*dr$ (same)

u_new : $u[0][i]*(1-dt*grad) + dtnu*Lapl$

This is quite significant for u_new (especially since it runs many times), factorising reduces the number of times the u array is accessed.

Branches and passing through data:

There are many if, else if statements within a nested loop – very expensive.

The only condition that needs to be within the nested loop is when both $ix, iy \neq 0$ or $Nx-1, Ny-1$.

All other conditions can be taken out of the inner loop, and some can be taken out of the outer loop.

Outside the outer loop are the following conditions:

- $ix, iy \neq 0$ (bottom left)
- $ix \neq 0, iy \neq Ny-1$
- $iy \neq 0, ix \neq Nx-1$
- $ix \neq Nx-1, iy \neq Ny-1$

These 4 conditions are exact.

Within the outer loop remain the following conditions:

David Degirmenci
u1624423

- $ix==0$ boundary (away from $iy==Ny-1,0$ boundaries)
- $iy==0$ boundary (away from $ix==Nx-1,0$ boundaries)
- $ix==Nx-1$ boundary (away from $iy==Ny-1,0$ boundaries)
- $iy==Ny-1$ boundary (away from $ix==Nx-1,0$ boundaries)

This is a large time save.

Optimising memory access and patterns for efficient use of cache:

In C, 2d arrays are stored in memory with row major ordering, so ordering accesses of arrays should be first in order of the first index then the second index.

This has been done for almost every array access. This is limited by multiplication and factorisation in general, but it is more expensive to have extra multiplication.

Further Optimisations:

Currently, each condition's code looks something like this:

```
for (i=1;i<Nx-1;i++) {  
  
    ///points away from boundary///  
    for (iy=1;iy<Ny-1;iy++) {  
        ///x  
        Lap1 = (u[i-1][iy] - 2.0*u[i][iy] + u[i+1][iy])*dx2r;  
        grad = (-u[i-1][iy] + u[i+1][iy])*dxr;  
        ///y  
        Lap1 += (u[i][iy-1] - 2.0*u[i][iy] + u[i][iy+1])*dy2r;  
        grad += (-u[i][iy-1] + u[i][iy+1])*dyr;  
  
        ///new value of u at this point  
        u_new[i][iy] = u[i][iy]*(1-dt*grad) + dt*nu*Lap1;  
    }  
    ///END///  
}
```

(this condition is within a nested loop, the others are not within the inner loop, and some are not even within the outer loop, although all are within the istep loop)

Both $dx, dy=1.0$, so $dx2r=dy2r=1.0$, and $dxr=dyr=0.5$. This means that these variables don't need to be used in the arithmetic. In fact

By setting Lap1, grad equal to the sum rather than adding on afterwards, we lose readability but we can reorder the array accesses for more efficient cache use. We can then place these into grad and Lap1 in u_new, so Lap1 and grad don't have to be used, shown here:

```
///points away from boundary///  
for (iy=1;iy<Ny-1;iy++) {  
    u_new[i][iy] = u[i][iy]*(1-dt*(-u[i-1][iy] - u[i][iy-1] + u[i][iy+1] + u[i+1][iy])*0.5) + dt*nu*(u[i-1][iy] + u[i][iy-1] - 4.0*u[i][iy] + u[i][iy+1] + u[i+1][iy]);  
}  
///END///
```

This entire optimisation is of negligible difference for each iteration, but the code above runs $10000 \times 254 \times 254 = 645,160,000$ times, so it adds up to up to a few seconds without any optimisation flags. I have done this similarly with the rest of the code, which saves a tiny amount of time.

Additionally, the way that u_new was being copied into u was very inefficient, with a nested loop (running a copy 650+million times). Ideally a simple `memcpy(u, u_new, Nx*Ny*sizeof(double))` could

David Degirmenci
u1624423

be used, but the way memory is allocated to these arrays means that is not possible. Instead slices of the array are iterated through and memcopied from `u_new` to `u`.

The snapshots of grid to file involves a branch, which runs 10000 times, but makes a negligible difference to the time, so this can be removed or kept (obviously kept if user wants to see and for there to be no compiler warnings).

Setting the optimisation compiler flags `-O1`, `-O2`, `-O3` reduces the time of the final code from on average 9.93 seconds to (average running on `stan1`):

- `-O1`: 3.11s
- `-O2`: 2.42s
- `-O3`: 1.65s

These all produce the same output as without any optimisation flags.

On the original code it reduces the time from 47.14s to (average running on `stan1`):

- `-O1`: 6.98s
- `-O2`: 5.47s
- `-O3`: 5.42s

(On my personal machine, with the `-O3` flag, I have managed to get the time down to less than 1 second)